


WORD2VEC AND WORD EMBEDDING VISUALIZATIONS


Brian Mayer

Research Scientist

Sanghani Center @ Virginia Tech




Demo Link: <https://github.com/mayer22/WordEmbeddings>



INTRODUCTION



WHO AM I?




Bachelor & Masters in
Industrial & Systems
Engineering

Logistics, Manufacturing, &
Performance Analysis

Transportation Research

Research Management,
Transportation, Intelligence,
and Forecasting Research

WHO AM I?



TODAY'S ASSUMPTIONS

- We code in Python
- We understand basic Natural Language Processing (NLP) techniques like:
 - Bag of Words (TF/TFIDF)
 - Basic sentiment analysis
- We know how neural networks work
- We know about dimension reduction

REVIEW — TEXT ANALYTICS

- **Bag of Words (BoW)** is an NLP technique that:
 - Uses words as variables
 - Ignores order, context, and meaning
 - Examples
 - Term Frequency (TF)
 - Create a Term-by-Document Matrix
 - Can highlight unimportant words because they appear a lot
 - TF-Inverse Document Frequency (TFIDF)
 - Evaluates the importance of a word to a document
 - Boost if a word is common to a document but penalize if it is also common to the corpus
 - Data representations become very large in large corpuses

What is an n -gram?

A group of n words

1-gram: “cat”, “dog”, “truck”, “bike”

2-gram (bi-gram): “fire truck”,
“yellow dog”, “my car”

3-gram (tri-gram): “ride my bike”,
“walk the dog”

4-gram: “go for a spin”,
“take a long walk”

... and so on (oops, that was a trigram)



REVIEW – TEXT ANALYTICS

- **Semantic Analysis** is an NLP technique that:
 - Represents meaning and considers words in context
 - Requires an emotional framework (labor intensive and requires many assumptions)
 - Requires many highly impactful parameters (e.g., valence shifters and window)
 - i.e., user specification





Figure 4.1 Plutchik's wheel of emotion with eight primary emotional states.

HOW DO WE DO BOTH?

Leverage the simplicity of counting words/appearance



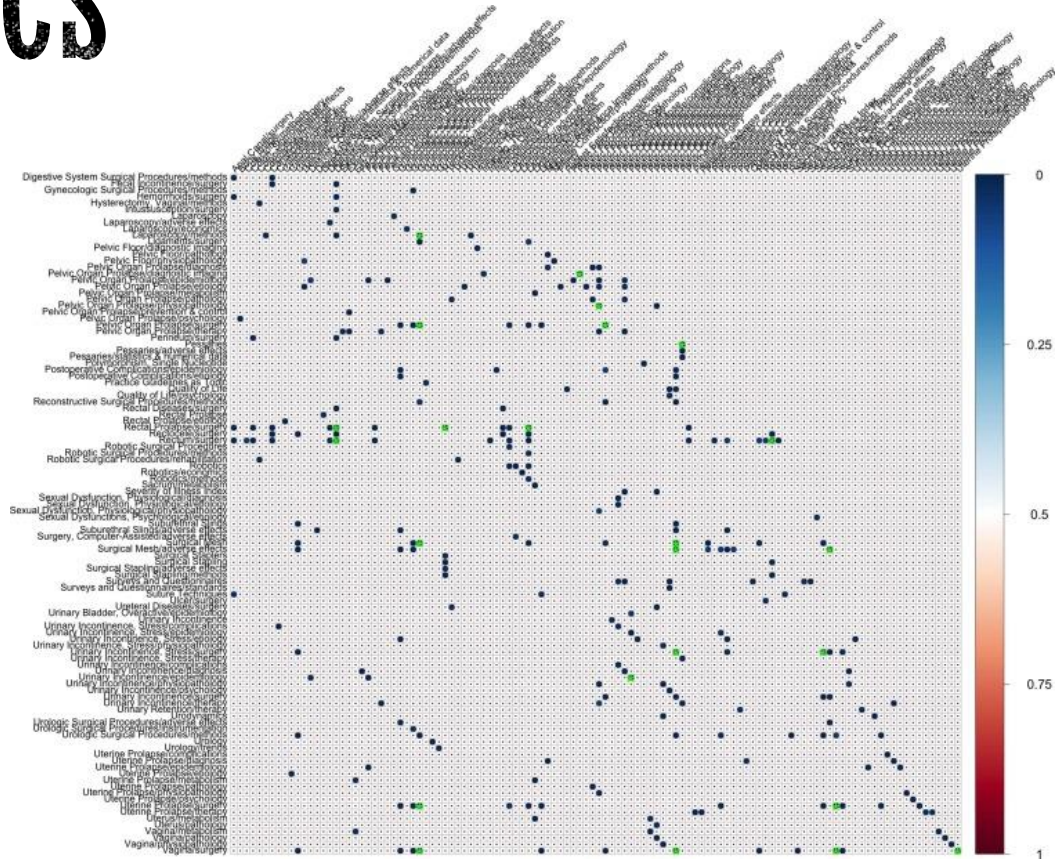
&

Consider the context of words



REVIEW — TEXT ANALYTICS

- A **Co-occurrence Matrix** is an NLP technique similar to Term-Frequency analysis that:
 - Counts how often different words appear together (within some window) in a document/corpus
 - i.e., adds context
 - Each word is represented as a vector of its co-occurrence frequencies with other words
 - Even larger than TF/TFIDF (Oh No!)



Zhou, Xiaobei, et al. "A probabilistic model for co-occurrence analysis in bibliometrics." *Journal of Biomedical Informatics* 128 (2022): 104047.



HOW CAN WE BE SMARTER?


- USE EMBEDDINGS => turn words into manageable vectors

? Why would we want to represent words as vectors?

1. Depict the meaning of a word (or document as a combination of words) quantitatively
 2. Compare words/documents based on location and relative distance
- Computational Linguistics and Social Science

- There are several different word embedding methods

? Are these outdated now?



word2vec

Co-Occurrence Matrix

GloVe

TFIDF

LARGE LANGUAGE MODELS (LLMs)

LLMs don't really compare words/docs BUT
ChatGPT evolved from word2vec


- We'll learn that embeddings are based on predictions (just like LLMs)
- 2018 ELMO¹ and BERT²: dynamic word embeddings
 - Embeddings change based on context
- 2018 Transformer Architecture³

2018 GPT1 

**So if we want to understand LLMs
let's go back to the start of all this**

1. Neumann, M. P. M., et al. "Deep contextualized word representations." *arXiv preprint arXiv:1802.05365* (2018).
2. Devlin, Jacob. "Bert: Pre-training of deep bidirectional transformers for language understanding." *arXiv preprint arXiv:1810.04805* (2018).
3. Vaswani, A. "Attention is all you need." *Advances in Neural Information Processing Systems* (2017).

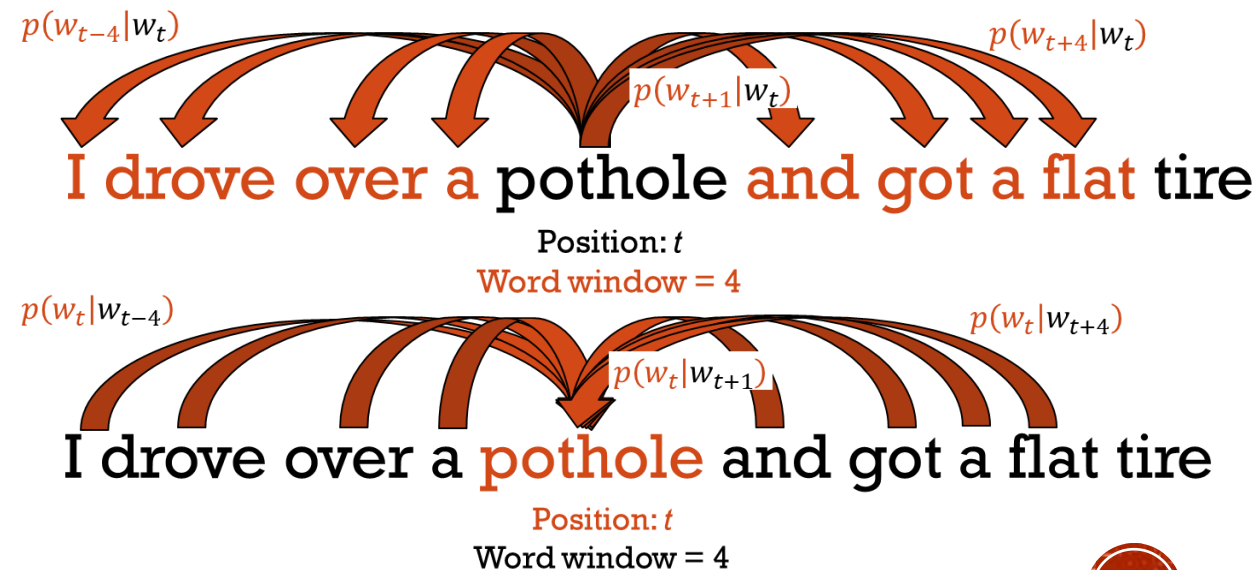





WORD2VEC

WORD2VEC (PREDICTION-BASED) EMBEDDINGS

- More advanced NLP method using neural networks
 - We actually use the weights not the output of the network
- Therefore, reliable results require a **significant amount of training data**
- Developed by researchers at Google
- Skip-Gram⁴
 - Trained to predict words before and after the current word
- Continuous Bag of Words (CBOW)⁵
 - Trained to predict a word based on context (nearby, i.e., related, words)


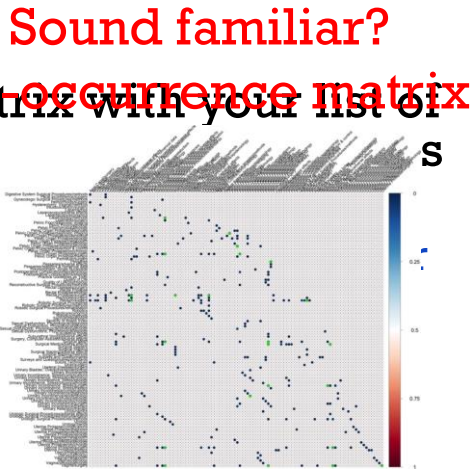


4. Mikolov, Tomas, et al. "Distributed representations of words and phrases and their compositionality." *Advances in neural information processing systems* 26 (2013).

5. Mikolov, Tomas. "Efficient estimation of word representations in vector space." *arXiv preprint arXiv:1301.3781* (2013).

LET'S START WITH SKIP-GRAM

- Find every word in your vocabulary/corpus
- Count how often words occur near each other
- Create an identity matrix with your list of words as labels for the (one-hot encoding)
 - a) Each vector in the matrix is a word
 - b) The output for training vectors matched with the every word we find near
- Create a NN with 1 hidden layer and N neurons (will be our # of parameters)
- Then train a softmax regression classifier to determine the output probabilities and hidden layer weights
 - Sum of all V probabilities per neuron will = 1 (per softmax rules)




MODEL EVALUATION

Example:

5 word corpus

with 3 features (weights)



Your hidden layer can have any number of features.

- 100 features is common
 - not too big (i.e., manageable)
 - not too small (i.e., still descriptive)
- Some people use smaller sizes to
 - save space
 - reduce computation
 - visualize in 2D

SUMMARY OF SKIP-GRAM NETWORK

- Training objective:

Maximize the probability of predicting nearby words correctly

- Takes a large input vector (V long), compresses it down to a small dense vector (N long), then outputs probability distribution for the target word
- The rows of the hidden layer weight matrix W_1 (columns of W_2) are the word vector.
- Last layer is always softmax (w/categorical cross-entropy)

SOFTMAX REGRESSION

$$p_i = \frac{e^{z_i}}{\sum_{j=1}^n e^{z_j}}$$

Vector $z = [z_1, z_2, \dots, z_V]$ (rows of W_1 , columns of W_2)

z_i represents the evaluated score associated with class i

p_i = probability of each class i

- Prior to the application of softmax some vector components could
 - be negative
 - greater than one
 - might not sum to 1
- But after applying the softmax function
 - each component will be in the interval (0,1)
 - the components will add up to 1
 - they can be interpreted as probabilities

UNDERSTANDING

If words frequently appear together (i.e., have similar contexts)

Our r **Hold that thought** words

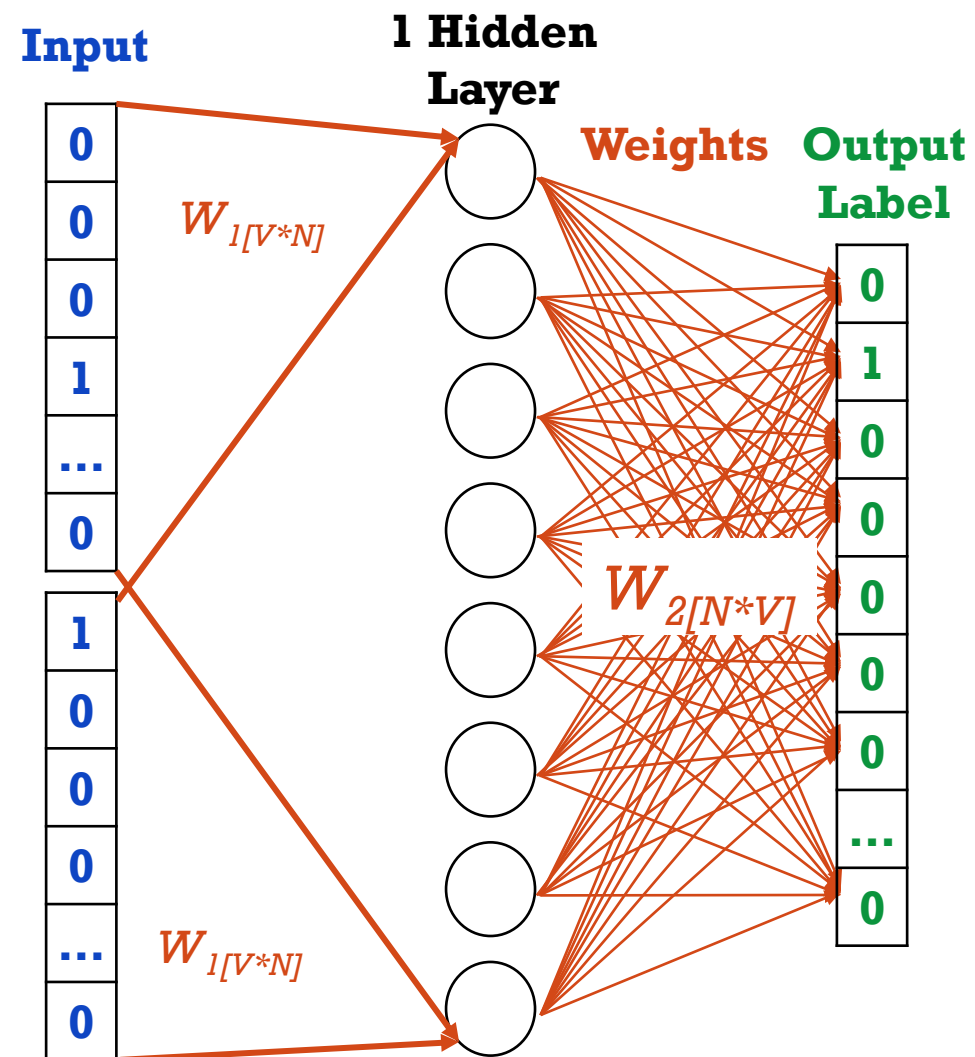
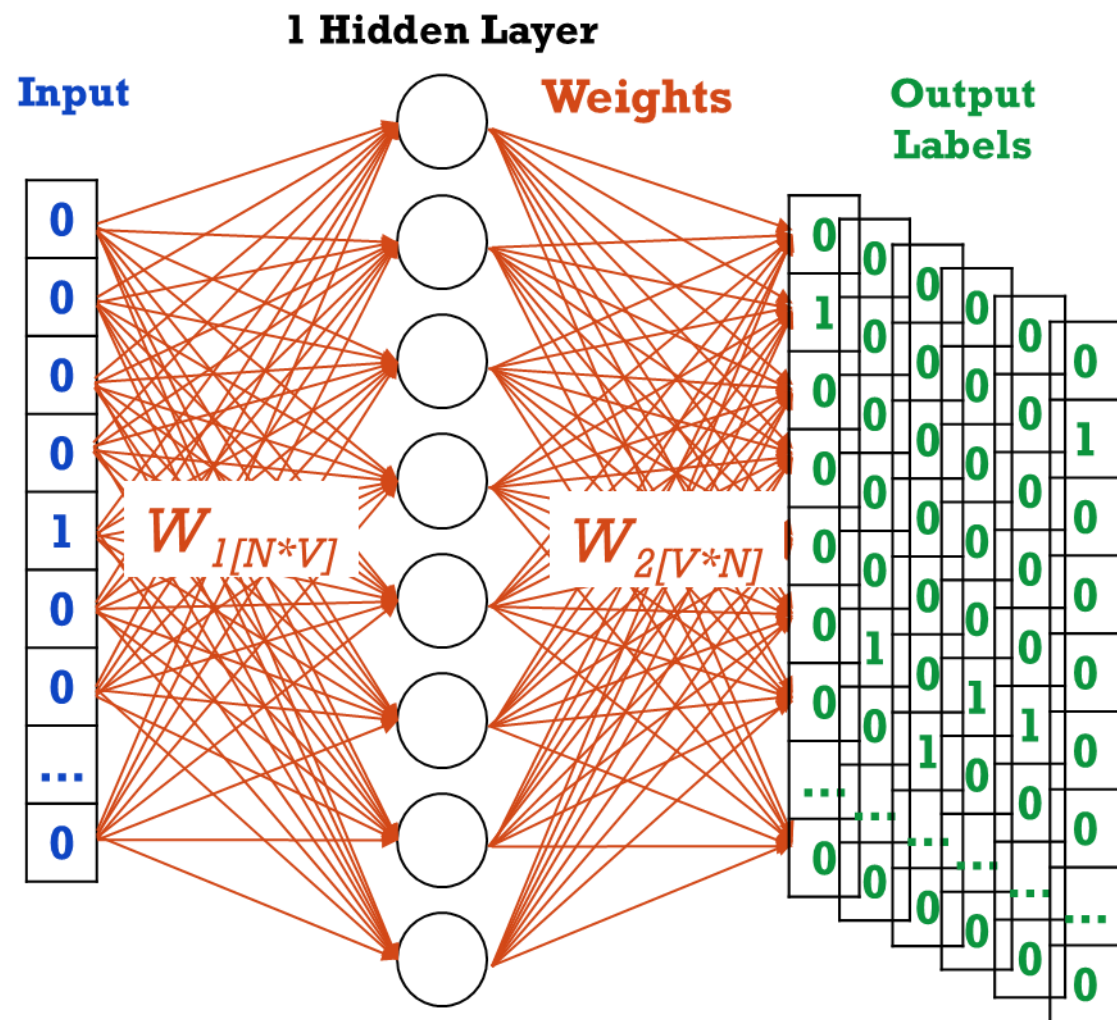
We'll come back to this

(i.e., re N is the number of weights/features)

SKIP-GRAM

VS

CBOW




SKIP-GRAM VS CBOW

■ Skip-Gram

- Better represents rare words and larger datasets
 - Considers each word pair
- **Slower**
- Better at representing *semantic meaning*


■ Continuous Bag of Words (CBOW)

- Common words are more accurate because they are smoothed over the entire dataset
- Rare words get “lost”/“forgotten”
- **Faster**: only needs to predict a single target word given a set of context words
- Better at representing *syntactic relationships*



WORD2VEC DEMO

<https://github.com/mayer22/WordEmbeddings>



WORD EMBEDDING VISUALIZATIONS



REMEMBER

If words frequently appear together (i.e., have similar contexts)

THEN

Our model needs to output similar results for the 2 words

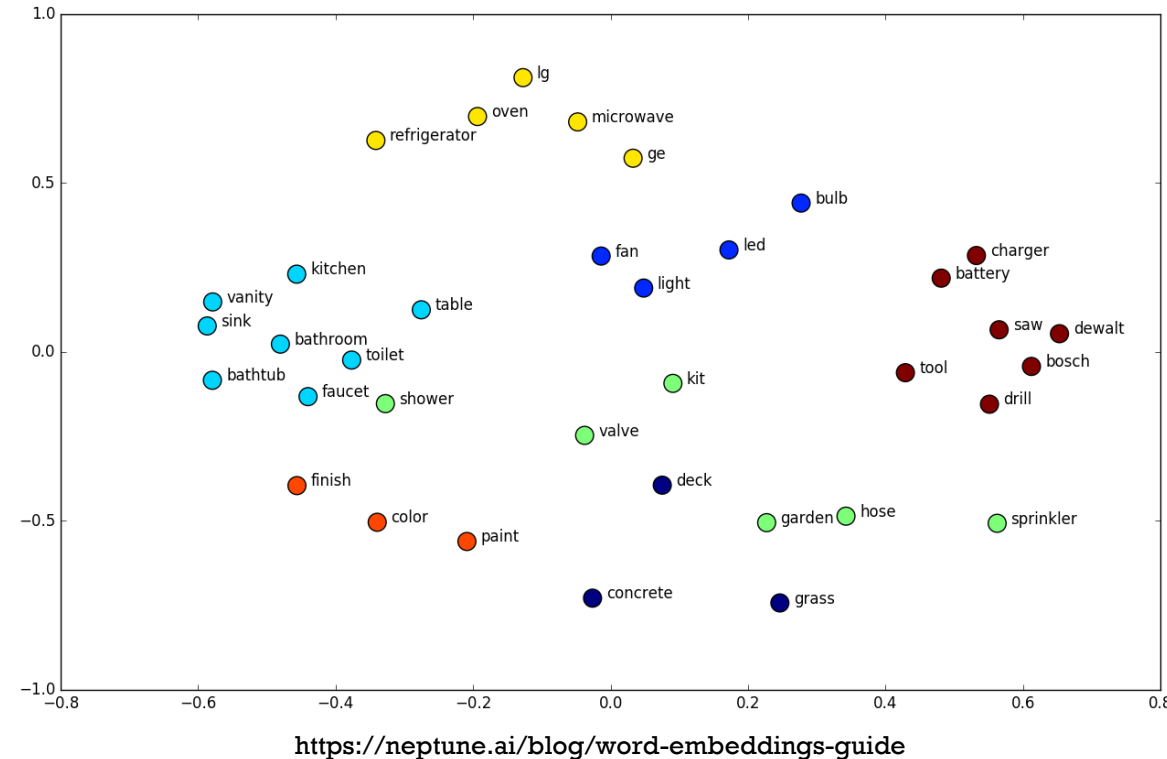
THEREFORE

Their word vectors will be similar

(i.e., relationally close in an n dimensional space where n is the number of weights/features)


COMPARE WORD VECTORS

- We can measure high-dimensional distances to:
 - find similar words
 - determine which word is more similar
 - find the opposite of a word (Italy?)
- We can visualize words on a scatter-plot
 - If you used more than 2 or 3 weights/parameters you will need to reduce them using these common methods:
 - Multidimensional Scaling (MDS)
 - Principal Component Analysis (PCA)
 - t-distributed stochastic neighbor embedding (t-SNE)
 - Uniform Manifold Approximation and Projection (UMAP)




COMPARE WORD VECTORS


Training Data:
Articles from
Google News
and classical
literary works
by Leo Tolstoy



<https://towardsdatascience.com/google-news-and-leo-tolstoy-visualizing-word2vec-word-embeddings-with-t-sne-11558d8bd4d>




**LET'S GO BACK TO OUR
DEMOS**




WORD EMBEDDING ANALYSIS

COMPARING WORD VECTORS


- Word analogies can often be solved with vector arithmetic⁶



Male-Female



Verb tense



Country-Capital

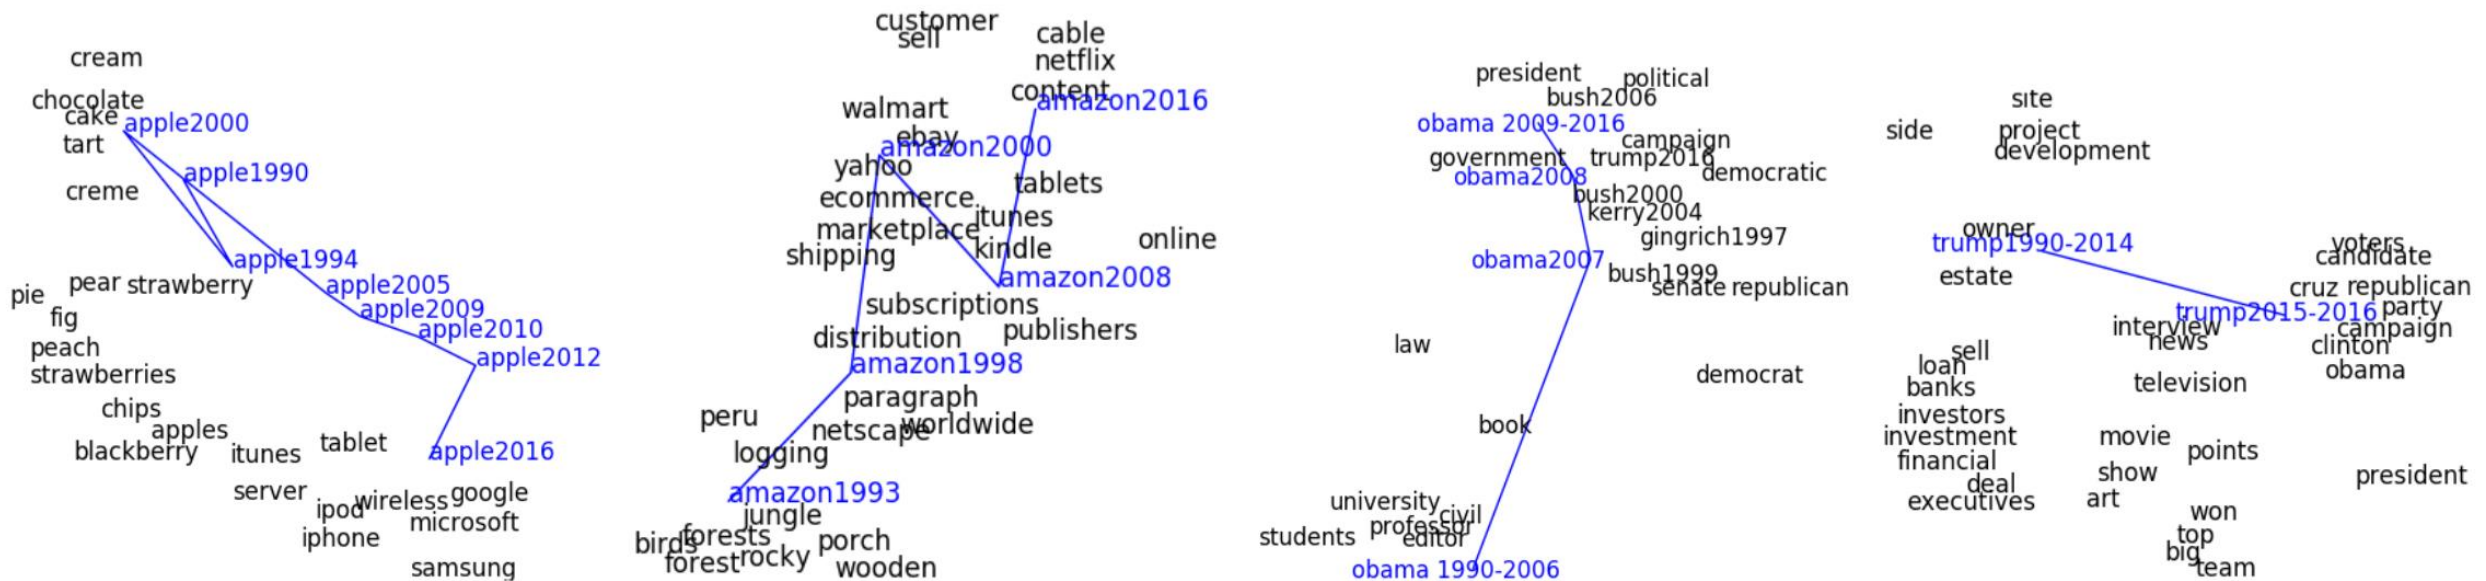
<https://www.analyticsvidhya.com/blog/2017/06/word-embeddings-count-word2veec/>

TEMPORAL WORD EMBEDDINGS

Word contexts change over time.

How do we view changes?


- Create temporal embeddings:^{7,8}



Global embedding space with chosen word vectors by time slice

TEMPORAL WORD EMBEDDINGS

- We can also train a model that updates the embeddings for all of our words every time slice
 - **But embeddings are not deterministic**
 - Embedding spaces can shift
 - Same word/same meaning => new location
 - **Alignment needed!**
- Sol: Force word embeddings in different time periods to be similar (NN training)^{9,10}
 - **Assumes** that the majority of the words do not change their meaning over time




9. Hamilton, William L., Jure Leskovec, and Dan Jurafsky. "Diachronic word embeddings reveal statistical laws of semantic change." *arXiv preprint arXiv:1605.09096* (2016).

10. Yao, Zijun, et al. "Dynamic word embeddings for evolving semantic discovery." *Proceedings of the eleventh acm international conference on web search and data mining*. 2018.

TEMPORAL WORD EMBEDDINGS WITH A COMPASS (TWEC)¹¹

Also, Compass-Aligned Distributional Embeddings (CADE)¹²

- Train compass using entire corpus
- Use embedding from Timestep 1 as a “compass” for other time slices
- **Assumes** that words that move, appear in the context of other words that move (more accurately accounts for shifted words)





11. Di Carlo, Valerio, Federico Bianchi, and Matteo Palmonari. "Training temporal word embeddings with a compass." Proceedings of the AAAI conference on artificial intelligence. Vol. 33. No. 01. 2019.

12. Bianchi, Federico, et al. "Compass-aligned distributional embeddings for studying semantic differences across corpora." arXiv preprint arXiv:2004.06519 (2020).

NEXT CLASS

Document Embeddings (Doc2Vec)
to Compare Documents





NOW LET'S EXPLORE WORD EMBEDDINGS

Tool: <https://www.cs.cmu.edu/~dst/WordEmbeddingDemo/index.html>

Tutorial: <https://www.cs.cmu.edu/~dst/WordEmbeddingDemo/tutorial.html>

Experiment: <https://www.cs.cmu.edu/~dst/WordEmbeddingDemo/experiments.html>

EXPLORATION ACTIVITY

(OR GO BACK TO CODE)

- Tool: <https://www.cs.cmu.edu/~dst/WordEmbeddingDemo/index.html>
- **Rotate**: Click and drag in the point plot to rotate the view
- **Zoom**: Use the scroll wheel to zoom in or out.
- **Pan**: Hold down the control key and click and drag to pan the view.
- **Word details**: Hovering over a word shows the closest 10 words
- **Activate**: Click on a word in the plot
 - Shows dot product/cosine similarity of the selected words to the words in the vector slots
 - Add it to the vector display on the right by clicking one of the slots
- Tutorial:
<https://www.cs.cmu.edu/~dst/WordEmbeddingDemo/tutorial.html>
- Experiment:
<https://www.cs.cmu.edu/~dst/WordEmbeddingDemo/experiments.html>

DISCUSS WITH YOUR NEIGHBORS

- Where would you put the words "adult", "child", "infant", or "grandfather"?
- See if you can perform some analogy arithmetic
 - Calculate the distances/directions between words for 2 sets of terms that make up an analogy and see if it is the same
- See if you can find terms that aren't similar but are similar in one parameter. Are there other terms similar/different on that parameter? What might that parameter be?
- Find another meaningful semantic dimension.
- <https://www.cs.cmu.edu/~dst/WordEmbeddingDemo/index.html>

QUESTIONS?

